Exercise 1

- Convert 8-bit data to $[0, 1]$ range and add per-pixel Gaussian noise $\epsilon \sim \mathcal{N}(0, 1/100)$.

- Model image patches as Gaussian

$$p_\theta(x) = \frac{1}{Z} e^{-\frac{1}{2}(x-\mu)^T \Lambda (x-\mu)}.$$

- Subtract the empirical mean, leading to

$$p_\theta(x) = \frac{1}{Z} e^{-\frac{1}{2} x^T \Lambda x}.$$

- Λ has a 2D Laplacian structure: 4-connected neighbouring pixels are correlated.

- Estimate Λ via

  - NCE (explain your choice of $p_n(x')$).
  - cNCE (explain your choice of $p_n(x'|x)$) or score matching (coin flip).

- Use SGD (or RMSProp or ADAM) for gradient-based optimization

- Visualize samples from p_θ (with $A = \Lambda^{-1/2} = \Sigma^{1/2}$)

$$x' \leftarrow \mu + A\epsilon, \qquad \epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I}).$$

- NCE: How close is the estimate of $\log Z$ to $\frac{D}{2} \log(2\pi) + \frac{1}{2} \log |\Sigma|$? $(D = 28^2)$

- Bonus exercise: Rerun with 8-connected neighbourhood assumption

Noise-contrastive estimation (NCE)

NCE casts the estimation of distribution parameters (in our case, the precision matrix Λ_θ) as a supervised learning problem. It also jointly estimates the unknown partition function Z_θ, and it does all this by applying logarithmic PSR to estimate parameters of a binary random variable.

Start by considering a batch of MNIST training images $x_1, x_2, x_3, \ldots$, which we assume are drawn from an unknown data distribution p_d. Also consider a fully known noise distribution $p_n(x')$; we could try a multidimensional Gaussian $\mathcal{N}(0, \Sigma_n)$, maybe with diagonal Σ_n? Perhaps $\Sigma_n = \epsilon I$?

Next, randomly let $z \in \{0, 1\}$ decide whether to draw a sample from p_d (i.e., drawing a random training image x_i) or a sample from the noise distribution p_n. That is,

$$p_{d,n}(x|z=0) = p_d(x), \qquad p_{d,n}(x|z=1) = p_n(x).$$

Further suppose that $p_z(z=0) = \eta$ and $p_z(z=1) = 1 - \eta$, where $\eta \in [0, 1]$ is a hyperparameter. Bayes rule combined with some formal symbol pushing yields the posterior distribution

$$\boxed{p_{d,n}(z|x) = \frac{(1-z)p_d(x) + \nu z p_n(x)}{p_d(x) + \nu p_n(x)}, \qquad \left(\eta = \frac{1}{1+\nu}\right)}$$

Given a sample x, we can thus calculate the probability $p_{d,n}(0|x)$ that x was drawn from the data distribution p_d, and the probability $p_{d,n}(1|x) = 1 - p_{d,n}(0|x)$ that x was drawn from the noise distribution p_n.

Posterior induced by true data distribution p_d:

$$p_{d,n}(z|x) = \frac{(1-z)p_d(x) + \nu z p_n(x)}{p_d(x) + \nu p_n(x)}$$

Posterior induced by model distribution p_θ:

$$p_{\theta,n}(z|x) = \frac{(1-z)p_\theta(x) + \nu z p_n(x)}{p_\theta(x) + \nu p_n(x)}$$

NCE uses logarithmic PSR to align these two posterior distributions. In order to do this, we need to form a set of training data $\{(x_i, z_i)\}$ where $z_i \sim p(z)$ and

$$x_i \sim \begin{cases} p_d(x) & \text{if } z = 0 \\ p_n(x) & \text{if } z = 1 \end{cases}$$

This should be very easy to implement in code:

1. Draw $z \in \{0, 1\}$ many times, from a Bernoulli distribution with probability $p(z = 0) = \eta$. Let N be the number of times you drew $z = 0$ and let M be the number of times you drew $z = 1$. (You can probably draw all these $N + M$ times using a single command, in one single line, you don't need a for-loop.)

2. For each time you drew $z = 0$, pick a random training example. For each time you drew $z = 1$, sample from the fully known noise distribution $p_n(x)$.

3. Store a list of the pairs (x_i, z_i).

In the next step, we will use primes and different subscripts to distinguish between training examples $x_i \sim p_d(x)$ and noise samples $x'_j \sim p_n(x)$. Note that expressions such as $p_d(x'_j)$ and $p_n(x_i)$ do make sense; we can always evaluate the probablity $p_n(x_i)$ of drawing the training example x_i from the noise distribution p_n and vice versa.

The goal now is to minimize the NCE objective

$$J(\theta) \propto \mathbb{E}_{x \sim p_d} \left[\log \frac{p_\theta(x)}{p_\theta(x) + \nu p_n(x)} \right] + \nu \mathbb{E}_{x \sim p_n} \left[\log \frac{\nu p_n(x)}{p_\theta(x) + \nu p_n(x)} \right]$$

$$\approx \frac{1}{N} \sum_{i=1}^{N} \log \frac{p_\theta(x_i)}{p_\theta(x_i) + \nu p_n(x_i)} + \frac{\nu}{M} \sum_{j=1}^{M} \log \frac{\nu p_n(x'_j)}{p_\theta(x'_j) + \nu p_n(x'_j)},$$

with $x_i \sim p_d$ and $x'_j \sim p_n$. Now suppose that

$$p_\theta(x) = \frac{1}{Z_\theta} e^{-\frac{1}{2}(x-\mu)^T \Lambda (x-\mu)} =: \frac{1}{Z_\theta} e^{f_\theta(x)},$$

$$p_n(x) = \frac{1}{Z_n} e^{-\frac{1}{2}(x-\mu_n)^T \Lambda_n (x-\mu_n)} =: \frac{1}{Z_n} e^{f_n(x)}.$$

Then we get

$$\log p_\theta(x) = f_\theta(x) - \log Z_\theta = -\frac{1}{2}(x-\mu)^T \Lambda(x-\mu) - \log Z_\theta,$$

$$\log p_n(x) = f_n(x) - \log Z_n = -\frac{1}{2}(x-\mu_n)^T \Lambda_n(x-\mu_n) - \log Z_n,$$

hence[1]

$$J(\theta) \approx \frac{1}{N} \sum_{i=1}^{N} \left(f_\theta(x_i) - \log Z_\theta - \log(p_\theta(x_i) + \nu p_n(x_i)) \right) +$$

$$+ \frac{\nu}{M} \sum_{j=1}^{M} \left(f_n(x'_j) - \log Z_n - \log\left(p_\theta(x'_j) + \nu p_n(x'_j)\right) + \log \nu \right)$$

$$= \frac{1}{N} \sum_{i=1}^{N} f_\theta(x_i) + \frac{\nu}{M} \sum_{j=1}^{M} f_n(x'_j) - \log Z_\theta - \nu \log Z_n + \nu \log \nu -$$

$$- \frac{1}{N} \sum_{i=1}^{N} \log(p_\theta(x_i) + \nu p_n(x_i)) - \frac{\nu}{M} \sum_{j=1}^{M} \log\left(p_\theta(x'_j) + \nu p_n(x'_j)\right)$$

We can also rewrite the last terms by noting that

$$\log(p_\theta(x) + \nu p_n(x)) = \log\left(\frac{1}{Z_\theta} e^{f_\theta(x)} + \frac{\nu}{Z_n} e^{f_n(x)} \right) \tag{1}$$

$$= \log\left(\left(\frac{Z_n}{Z_\theta} e^{f_\theta(x) - f_n(x)} + \nu \right) \left(\frac{1}{Z_n} e^{f_n(x)} \right) \right)$$

$$= \log\left(\frac{Z_n}{Z_\theta} e^{f_\theta(x) - f_n(x)} + \nu \right) + f_n(x) - \log Z_n,$$

---

[1] I write $J(\theta) \approx$, but what I mean is that $J(\theta)$ is proportional to something which is approximately equal to ...

hence

$$J(\theta) \approx \frac{1}{N} \sum_{i=1}^{N} f_\theta(x_i) + \frac{\nu}{M} \sum_{j=1}^{M} f_n(x_j') - \log Z_\theta - \nu \log Z_n + \nu \log \nu -$$

$$- \frac{1}{N} \sum_{i=1}^{N} \log(p_\theta(x_i) + \nu p_n(x_i)) - \frac{\nu}{M} \sum_{j=1}^{M} \log\big(p_\theta(x_j') + \nu p_n(x_j')\big)$$

$$= \frac{1}{N} \sum_{i=1}^{N} f_\theta(x_i) + \frac{\nu}{M} \sum_{j=1}^{M} f_n(x_j') - \log Z_\theta - \nu \log Z_n + \nu \log \nu -$$

$$- \frac{1}{N} \sum_{i=1}^{N} \left(\log \left(\frac{Z_n}{Z_\theta} e^{f_\theta(x_i) - f_n(x_i)} + \nu \right) + f_n(x_i) - \log Z_n \right) -$$

$$- \frac{\nu}{M} \sum_{j=1}^{M} \left(\log \left(\frac{Z_n}{Z_\theta} e^{f_\theta(x_j') - f_n(x_j')} + \nu \right) + f_n(x_j') - \log Z_n \right)$$

$$= \frac{1}{N} \sum_{i=1}^{N} \left[f_\theta(x_i) - f_n(x_i) \right] + \log(Z_n/Z_\theta) + \nu \log \nu -$$

$$- \frac{1}{N} \sum_{i=1}^{N} \log \left(\frac{Z_n}{Z_\theta} e^{f_\theta(x_i) - f_n(x_i)} + \nu \right) - \frac{\nu}{M} \sum_{j=1}^{M} \log \left(\frac{Z_n}{Z_\theta} e^{f_\theta(x_j') - f_n(x_j')} + \nu \right)$$

We can also rewrite this in another way, by factorizing out p_θ instead of p_n in equation (1). Then we should get something similar to

$$J(\theta) \approx \frac{\nu}{M} \sum_{j=1}^{M} \left[f_n(x_j') - f_\theta(x_j') \right] + \nu \log(\nu Z_\theta / Z_n) -$$

$$- \frac{1}{N} \sum_{i=1}^{N} \log \left(\nu \frac{Z_\theta}{Z_n} e^{f_n(x_i) - f_\theta(x_i)} + 1 \right) - \frac{\nu}{M} \sum_{j=1}^{M} \log \left(\nu \frac{Z_\theta}{Z_n} e^{f_n(x_j') - f_\theta(x_j')} + 1 \right)$$

I spent a very small amount of time trying to see if these sums of logarithms simplify, using the additive law of logarithms, but they don't seem to. I also tried calculating the average of these two approximations of $J(\theta)$ to see if that would simplify the logarithms, but that didn't seem to give me anything either.

Now suppose that we have normalized so that $p_\theta \sim \mathcal{N}(0, \Sigma_\theta)$ and $p_n \sim \mathcal{N}(0, \Sigma_n)$ with $\Lambda = \Sigma^{-1}$. Using the identity $\det(\Sigma) = 1/\det(\Lambda)$, the above approximation of $J(\theta)$ becomes

$$
\boxed{
\begin{aligned}
J(\theta) \approx &\frac{\nu}{M}\left(\sum_{j=1}^{M} x_j'^{\,T}(\Lambda_n - \Lambda_\theta)x_j'\right) + \frac{\nu}{2}\log\left(\nu^2 \frac{\det \Lambda_n}{\det \Lambda_\theta}\right) - \\
&- \frac{1}{N}\sum_{i=1}^{N}\log\left(\nu\sqrt{\frac{\det \Lambda_n}{\det \Lambda_\theta}}\exp\left(x_i^{\,T}(\Lambda_n - \Lambda_\theta)x_i\right) + 1\right) - \\
&- \frac{\nu}{M}\sum_{j=1}^{M}\log\left(\nu\sqrt{\frac{\det \Lambda_n}{\det \Lambda_\theta}}\exp\left(x_j'^{\,T}(\Lambda_n - \Lambda_\theta)x_j'\right) + 1\right)
\end{aligned}
}
$$

This expression honestly doesn't look too bad, we can definitely compute every component. Computing the gradient of this approximation is fairly straightforward (and can likely be automated with Gradient Tape, so the following computations may be unnecessary). Before computing the derivatives, recall that the derivative of a parameterized matrix is defined by computing the derivative of each component:

$$
\left(\frac{\partial \Lambda_\theta}{\partial \theta}\right)_{ij} := \frac{\partial (\Lambda_\theta)_{ij}}{\partial \theta},
$$

where we have pretended that θ is a single variable. In practice, θ is an array containing multiple parameters θ_k, and the corresponding expressions for derivatives with respect to θ_k are obtained by simply adding the subscript $_k$ at the appropriate places:

$$
\left(\frac{\partial \Lambda_\theta}{\partial \theta_k}\right)_{ij} = \frac{\partial \Lambda_{ij}}{\partial \theta_k}.
$$

For notational convenience we shall also suppress the subscript when computing $\frac{\partial J}{\partial \theta_k}$, because we can just put the subscript back afterwards. We find that

$$
\begin{aligned}
\frac{\partial J}{\partial \theta} \approx &- \frac{\nu}{M}\left(\sum_{j=1}^{M} x_j'^{\,T}\frac{\partial \Lambda_\theta}{\partial \theta}x_j'\right) + \frac{\nu}{2\det \Lambda_\theta}\frac{\partial \det \Lambda_\theta}{\partial \theta} - \\
&- \frac{1}{N}\sum_{i=1}^{N}\frac{\nu\sqrt{\det \Lambda_n}\left[-\frac{1}{2}(\det \Lambda_\theta)^{-3/2}\frac{\partial \det \Lambda_\theta}{\partial \theta} - (\det \Lambda_\theta)^{-1/2}\left(x_i^T\frac{\partial \Lambda_\theta}{\partial \theta}x_i\right)\right]\exp\left(x_i^T(\Lambda_n - \Lambda_\theta)x_i\right)}{\nu\sqrt{\frac{\det \Lambda_n}{\det \Lambda_\theta}}\exp\left(x_i^T(\Lambda_n - \Lambda_\theta)x_i\right) + 1} - \\
&- \frac{\nu}{M}\sum_{j=1}^{M}\frac{\nu\sqrt{\det \Lambda_n}\left[-\frac{1}{2}(\det \Lambda_\theta)^{-3/2}\frac{\partial \det \Lambda_\theta}{\partial \theta} - (\det \Lambda_\theta)^{-1/2}\left(x_j'^{\,T}\frac{\partial \Lambda_\theta}{\partial \theta}x_j'\right)\right]\exp\left(x_j'^{\,T}(\Lambda_n - \Lambda_\theta)x_j'\right)}{\nu\sqrt{\frac{\det \Lambda_n}{\det \Lambda_\theta}}\exp\left(x_j'^{\,T}(\Lambda_n - \Lambda_\theta)x_j'\right) + 1},
\end{aligned}
$$

which can be greatly simplified by cancelling factors from both numerators and denominators:

$$\frac{\partial J}{\partial \theta} \approx -\frac{\nu}{M}\left(\sum_{j=1}^{M} x_j'^T \frac{\partial \Lambda_\theta}{\partial \theta} x_j'\right) + \frac{\nu}{2\det\Lambda_\theta}\frac{\partial \det \Lambda_\theta}{\partial \theta} +$$

$$+ \frac{\nu}{N}\sum_{i=1}^{N} \frac{\frac{1}{2\det\Lambda_\theta}\frac{\partial \det \Lambda_\theta}{\partial \theta} + x_i^T \frac{\partial \Lambda_\theta}{\partial \theta} x_i}{\nu + \sqrt{\frac{\det \Lambda_\theta}{\det \Lambda_n}}\exp\left(-x_i^T(\Lambda_n - \Lambda_\theta)x_i\right)} +$$

$$+ \frac{\nu^2}{M}\sum_{j=1}^{M} \frac{\frac{1}{2\det\Lambda_\theta}\frac{\partial \det \Lambda_\theta}{\partial \theta} + x_j'^T \frac{\partial \Lambda_\theta}{\partial \theta} x_j'}{\nu + \sqrt{\frac{\det \Lambda_\theta}{\det \Lambda_n}}\exp\left(-x_j'^T(\Lambda_n - \Lambda_\theta)x_j'\right)}$$

This expression can be simplified further, using Jacobi's formula:

$$\frac{\partial \det \Lambda_\theta}{\partial \theta} = \mathrm{tr}\left(\mathrm{adj}(\Lambda_\theta)\frac{\partial \Lambda_\theta}{\partial \theta}\right),$$

where $\mathrm{adj}(\Lambda_\theta)$ is the adjugate of Λ_θ. Since Λ_θ is invertible, $\mathrm{adj}(\Lambda_\theta) = (\det \Lambda_\theta)\Lambda_\theta^{-1} = (\det \Lambda_\theta)\Sigma_\theta$, so after reintroducing the subscript k we get

$$\frac{\partial J(\theta)}{\partial \theta_k} \approx -\frac{\nu}{M}\left(\sum_{j=1}^{M} x_j'^T \frac{\partial \Lambda_\theta}{\partial \theta_k} x_j'\right) + \frac{\nu}{2}\mathrm{tr}\left(\Sigma_\theta \frac{\partial \Lambda_\theta}{\partial \theta_k}\right) +$$

$$+ \frac{\nu}{N}\sum_{i=1}^{N} \frac{\frac{1}{2}\mathrm{tr}\left(\Sigma_\theta \frac{\partial \Lambda_\theta}{\partial \theta_k}\right) + x_i^T \frac{\partial \Lambda_\theta}{\partial \theta_k} x_i}{\nu + \sqrt{\frac{\det \Lambda_\theta}{\det \Lambda_n}}\exp\left(-x_i^T(\Lambda_n - \Lambda_\theta)x_i\right)} +$$

$$+ \frac{\nu^2}{M}\sum_{j=1}^{M} \frac{\frac{1}{2}\mathrm{tr}\left(\Sigma_\theta \frac{\partial \Lambda_\theta}{\partial \theta_k}\right) + x_j'^T \frac{\partial \Lambda_\theta}{\partial \theta_k} x_j'}{\nu + \sqrt{\frac{\det \Lambda_\theta}{\det \Lambda_n}}\exp\left(-x_j'^T(\Lambda_n - \Lambda_\theta)x_j'\right)}$$

This expression is bulky, but if we can compute $\frac{\partial \Lambda_\theta}{\partial \theta_k}$ and $\Sigma_\theta = \Lambda_\theta^{-1}$ efficiently, then the whole derivative should be relatively efficient to compute. Indeed, the expression mainly contains determinants of known matrices, matrix products of known matrices, and scalar products. Of course, we need to compute this derivative for each component θ_k but we can save resources by only computing the k-independent weights

$$w(x) = \sqrt{\frac{\det \Lambda_n}{\det \Lambda_\theta}}\exp\left(x^T(\Lambda_n - \Lambda_\theta)x\right)$$

for $x \in \{x_1, \ldots, x_N, x_1', \ldots, x_M'\}$ once each time we update θ (that is, once per epoch). We can then reuse these pre-computed weights when computing each derivative $\frac{\partial J(\theta)}{\partial \theta_k}$, which then has the simpler expression

$$\frac{\partial J(\theta)}{\partial \theta_k} \approx \frac{\nu}{N} \sum_{i=1}^{N} \frac{w(x_i)}{\nu w(x_i) + 1} \left(x_i^T \frac{\partial \Lambda_\theta}{\partial \theta_k} x_i \right) -$$

$$- \frac{\nu}{M} \sum_{j=1}^{M} \left(\frac{w(x_j')}{\nu w(x_j') + 1} - 1 \right) \left({x_j'}^T \frac{\partial \Lambda_\theta}{\partial \theta_k} x_j' \right) +$$

$$+ \frac{1}{2} \operatorname{tr}\left(\Sigma_\theta \frac{\partial \Lambda_\theta}{\partial \theta_k} \right) \left(\nu + \frac{1}{N} \sum_{i=1}^{N} \frac{w(x_i)}{\nu w(x_i) + 1} + \frac{\nu}{M} \sum_{j=1}^{M} \frac{w(x_j')}{\nu w(x_j') + 1} \right)$$

Many parts of this expression are independent of k and can be pre-computed. Moreover, it would be nice if we could rewrite $\Sigma_\theta \frac{\partial \Lambda_\theta}{\partial \theta_k}$ as the derivative $\frac{\partial M_\theta}{\partial \theta_k}$ of some matrix M_θ, because then we could use linearity to rewrite the trace term as

$$\frac{\partial \operatorname{tr}(M)}{\partial \theta_k},$$

but I don't know if such a matrix M exists. It might also be computationally cheaper to undo our earlier application of Jacobi's formula and, instead of the trace, compute

$$\frac{1}{\det(\Lambda_\theta)} \frac{\partial \det(\Lambda_\theta)}{\partial \theta_k},$$

but I doubt it; the trace shold be cheaper as we need to compute $\frac{\partial \Lambda_\theta}{\partial \theta_k}$ anyway.

It might very well be possible to construct a tensor $\nabla_\theta \Lambda_\theta$, with matrix-valued components $\frac{\partial \Lambda_\theta}{\partial \theta_k}$, and compute $\nabla J(\theta)$ in one go. That depends on how capable Tensorflow / PyTorch / ... is.

Remark: If $\frac{\partial \Sigma_\theta}{\partial \theta_k}$ for some reason is easier to compute than $\frac{\partial \Lambda_\theta}{\partial \theta_k}$, then we can rewrite the above expression using the equalities

$$\frac{\partial \Lambda_\theta}{\partial \theta_k} = -\Lambda_\theta \frac{\partial \Sigma_\theta}{\partial \theta_k} \Lambda_\theta, \qquad \Sigma_\theta \frac{\partial \Lambda_\theta}{\partial \theta_k} = -\frac{\partial \Sigma_\theta}{\partial \theta_k} \Lambda_\theta.$$

Implementation steps for NCE:

Step 0. Load MNIST minibatch $\{x_1, x_2, x_3, \ldots\}$ which has been preprocessed, e.g. the empirical mean of the entire MNIST dataset should be subtracted from each image. Next, define the model distribution $p_\theta \sim \mathcal{N}(0, \Sigma_\theta)$ for some initial weights $\theta \in \mathbb{R}^K$, and a noise distribution $p_n \sim \mathcal{N}(0, \Sigma_n)$. Then choose a hyperparameter $\eta \in [0,1]$ for the distribution $p_z(z=0) = \eta$ and compute $\nu = \frac{1}{\eta} - 1$.

Step 1. Draw $z \sim p_z$ a predetermined number of times. (This number is a hyperparameter?)

- Each time $z = 0$, draw a random training example x_i from the minibatch.

- Each time $z = 1$, draw x_j' from the noise distribution p_n.

Let N be the number of times you drew $z = 0$, and M the number of times you drew $z = 1$.

Step 2. Use the samples obtained above to compute the weights

$$w(x) = \sqrt{\frac{\det \Lambda_n}{\det \Lambda_\theta}} \exp\left(x^T(\Lambda_n - \Lambda_\theta)x\right)$$

for all samples $x \in \{x_1, \ldots, x_N, x_1', \ldots, x_M'\}$.

Step 3. Compute the NCE objective[2]

$$\boxed{J(\theta) \approx \frac{\nu}{M} \sum_{j=1}^{M} \log\left(\frac{\nu w(x_j')}{\nu w(x_j') + 1}\right) - \frac{1}{N} \sum_{i=1}^{N} \log\left(\nu w(x_i) + 1\right)}$$

Step 4. Minimize the NCE objective, either using Gradient Tape or using the relation

$$\boxed{\begin{aligned}
\frac{\partial J(\theta)}{\partial \theta_k} &\approx \frac{\nu}{N} \sum_{i=1}^{N} \frac{w(x_i)}{\nu w(x_i) + 1} \left(x_i^T \frac{\partial \Lambda_\theta}{\partial \theta_k} x_i\right) - \\
&\quad - \frac{\nu}{M} \sum_{j=1}^{M} \left(\frac{w(x_j')}{\nu w(x_j') + 1} - 1\right) \left(x_j'^{\,T} \frac{\partial \Lambda_\theta}{\partial \theta_k} x_j'\right) + \\
&\quad + \frac{1}{2} \operatorname{tr}\left(\Sigma_\theta \frac{\partial \Lambda_\theta}{\partial \theta_k}\right) \left(\nu + \frac{1}{N} \sum_{i=1}^{N} \frac{w(x_i)}{\nu w(x_i) + 1} + \frac{\nu}{M} \sum_{j=1}^{M} \frac{w(x_j')}{\nu w(x_j') + 1}\right)
\end{aligned}}$$

for each parameter θ_k. It might be possible to construct a tensor $\nabla_\theta \Lambda_\theta$ so the full gradient $\nabla J(\theta)$ can be computed in one go. Regardless, I hope we can use Gradient Tape to compute each derivative $\frac{\partial \Lambda_\theta}{\partial \theta_k}$ because I don't know what to do otherwise.

---

[2]The approximation sign is not really correct. It is an approximation of an expectation *proportional to* $J(\theta)$.

Exercise 2

- Extract 50,000 random 28×28 natural image patches from a set of images
    - Holiday pictures, public domain pictures
- Convert patches to grayscale $\in [0,1]^{28 \times 28}$
- Compute a (constrained) Gaussian representing the data
    - Computing the empirical mean
    - Using the method from Exercise 1 to estimate a constrained precision matris Λ
    - Using the method of your choice (SM, NCE, cNCE)
    - We will optionally whiten the patches using this estimate in later steps
- Train a 2-layer DEM of the following form (using a method of your choice):

$$\log p_\theta(x) := -\frac{1}{2\sigma^2}\|x\|^2 + b^T x + \sum_{k=1}^{K} S\left(w_k^T g_\theta(x) + c_k\right)$$

$$g_\theta(x) = s(Vx) \quad \text{single layer sigmoid NN}$$

where (numerically robust expressions for) S and s are given by

$$S(u) = \begin{cases} \log(1+e^u) & u \leq 0 \\ u + \log\left(1 + e^{-u}\right) & u \geq 0 \end{cases}, \quad \text{and} \quad s(u) = \begin{cases} \dfrac{e^u}{1+e^u}, & u \leq 0 \\ \dfrac{1}{1+e^{-u}}, & u \geq 0 \end{cases}$$

As x and Vx are vector-valued, I assume that $s(Vx)$ is to be interpreted element-wise.

- Parameters: W, v, b, c e.g. $K = 64$, $V = \mathbb{R}^{64 \times 768}$.
- Choose $\sigma \in \{1, 0.1\}$
- Run on whitened and non-whitened data
- Visualize filters v_j (rows in V)
- Is your model able to distinguish between the following?
    - Hold-out natural patches
    - Samples generated by the fitted Gaussian ($\mathcal{N}(\mathbf{0}, \mathbf{1})$ in the whitened version)
    - MNIST digits
- Report mean and std of $\log p_\theta$
- Bonus: Generate modes of p_θ by maximizing $\log p_\theta(x)$
    - Inital x: Sample from fitted Gaussian

Contrastive divergence

We assume the model takes the shape

$$p_\theta(x, z) = \frac{e^{f_\theta(x,z)}}{Z(\theta)}, \quad \text{with marginal} \quad p_\theta(x) = \int \frac{e^{f_\theta(x,z)}}{Z(\theta)} \mathrm{d}z.$$

where the exponent in our case seems to be given by[3]

$$f_\theta(x, z) = \log p_\theta(x, z) = -\frac{1}{2\sigma^2}\|x\|^2 + b^T x + c^T z + z^T w g_\theta(x).$$

The presentation shows on page 47 that the gradient of Maximum Likelihood Estimation (MLE) for unnormalized latent variable models is given by

$$\nabla_\theta \log p_\theta(x) = \mathbb{E}_{z \sim p_\theta(\cdot|x)} \left[\nabla_\theta f_\theta(x, z) \right] - \mathbb{E}_{(x',z) \sim p_\theta} \left[\nabla_\theta f_\theta(x', z) \right],$$

where the first term can be estimated form samples:

$$\begin{cases} \mathbb{E}_{z \sim p_\theta(\cdot|x)} \left[\nabla_\theta f_\theta(x, z) \right] \approx \dfrac{1}{N} \sum_{i=1}^{N} \nabla_\theta f_\theta(x, z_i), \\ z_i \sim p_\theta(z|x) = \prod_k \mathrm{Ber}(z^k; s(w_k^T g_\theta(x) + c_k)) \end{cases}$$

(These equations combine pages 47 & 60. Note that ^{k} is an upper index over vector components, not an exponent.) The latent vectors

$$z_i = \begin{bmatrix} z_i^1 \\ \vdots \\ z_i^K \end{bmatrix}$$

are thus sampled by drawing each individual component $z_i^k = 1$ with probability $s(w_k^T g_\theta(x) + c_k)$ and $z_i^K = 0$ with probability $1 - s(w_k^T g_\theta(x) + c_k)$.

The second term is often harder, since it requires drawing samples from the joint distribution, but we can solve that problem in our case because we already know the marginalization with respect to z. The key is to first draw a random sample $x' \sim p_\theta(x)$ from the marginal distribution and then draw $z \sim p_\theta(z|x')$ as described above. Repeating this process M times yields a sample set $\{(x'_j, z_j)\}_{j=1}^{M}$ and we find that

$$\mathbb{E}_{(x',z) \sim p_\theta} \left[\nabla_\theta f_\theta(x', z) \right] \approx \frac{1}{M} \sum_{j=1}^{M} \nabla_\theta f_\theta(x'_j, z_j).$$

---

[3]The only possible difference between $f_\theta(x, z)$ and $\log p_\theta(x, z)$ is that the latter should include $-\log Z(\theta)$, unless Christopher just ignored that part. And I don't think any of the terms in the given expression for $\log p_\theta(x, z)$ is the term $-\log Z(\theta)$ in disguise, so it's probably safe to say that the model is unnormalized and $f_\theta(x, z) = \log p_\theta(x, z)$.

Computing derivatives and gradients

Let's compute gradients, based on the expression

$$\log p_\theta(x) = -\frac{1}{2\sigma^2}\|x\|^2 + b^T x + \sum_{k=1}^{K} S(w_k^T g_\theta(x) + c_k).$$

Hopefully this can be done automatically since most (all?) of these terms can be written in terms of Keras layers. For example, I guess we don't have to compute $\nabla_V \log p_\theta(x)$ as the layer $s(Vx)$ is just a dense layer without bias vector and with the sigmoid activation function.

In what follows, we set $u_k = w_k^T g_\theta(x) + c_k$ and use the relation

$$\frac{\partial S(u)}{\partial u} = \frac{\partial \log(1 + e^u)}{\partial u} = \frac{e^u}{1 + e^u} = s(u).$$

- Gradient with respect to b:
$$\nabla_b \log p_\theta(x) = x$$

- Gradient with respect to c:
$$\frac{\partial \log p_\theta(x)}{\partial c_k} = \frac{\partial S(u_k)}{\partial u_k}\frac{\partial u_k}{\partial c_k} = s(u_k),$$

  independently of k. Hence $\nabla_k \log p_\theta(x) = \begin{bmatrix} s(u_1) & \cdots & s(u_K) \end{bmatrix}^T$.

- Gradient with respect to each row w_k in the w matrix:
$$\frac{\partial \log p_\theta(x)}{\partial w_{kl}} = \frac{\partial S(u_k)}{\partial u_k}\frac{\partial u_k}{\partial w_{kl}} = s(u_k)(g_\theta(x))_l,$$

  hence $\nabla_{w_k} \log p_\theta(x) = \underbrace{s(u_k)}_{\text{scalar}}\underbrace{g_\theta(x)}_{\text{vector}}$.

- Derivatives with respect to V: Recalling that $g_\theta = s(Vx)$, we have
$$\frac{\partial \log p_\theta(x)}{\partial V_{ij}} = \sum_{k=1}^{K} \frac{\partial S(u_k)}{\partial u_k}\frac{\partial u_k}{\partial V_{ij}} = \sum_{k=1}^{K} s(u_k)\frac{\partial u_k}{\partial V_{ij}},$$

  where
$$\frac{\partial u_k}{\partial V_{ij}} =$$

Whitening images

Pixels in images are highly correlated; whitening aims to remove 2nd order correlations from images, so that the learned models can focus on higher-order correlations. Two kinds of whitening are expressed in the presentation: PCA and ZCA whitening. The two methods are similar:

1. Subtract empirical mean (I don't know what purpose $\mathbb{1}$ serves here)

$$X \leftarrow X - \frac{1}{N} X \mathbb{1}, \qquad \mathbf{X} = [x_1, \ldots, x_n]$$

2. Empirical covariance matrix and eigendecomposition:

$$C \leftarrow \frac{1}{N-1} X X^T, \qquad C = U \Lambda U^T.$$

3. Whitening step.

    - **PCA:** $\quad X \leftarrow \quad \Lambda^{-1/2} U^T X$
    - **ZCA:** $\quad X \leftarrow U^T \Lambda^{-1/2} U^T X \quad$ (Also rotates back to original space)

ZCA whitening enhances edges in images, while PCA whitening leads to unrecognizable images (as we don't rotate back to the original space).